



Machine Learning

From Turing to Deep Learning

Two approaches in AI

Good Old-fashion AI

- **Algorithms !**
- Design from human knowledge.
- Fast, thoughtful and less dependant on data.
- Centralized processes (analyse, maths, blueprints)
- Examples : GOFAI
- Less efficient

Data-driven AI

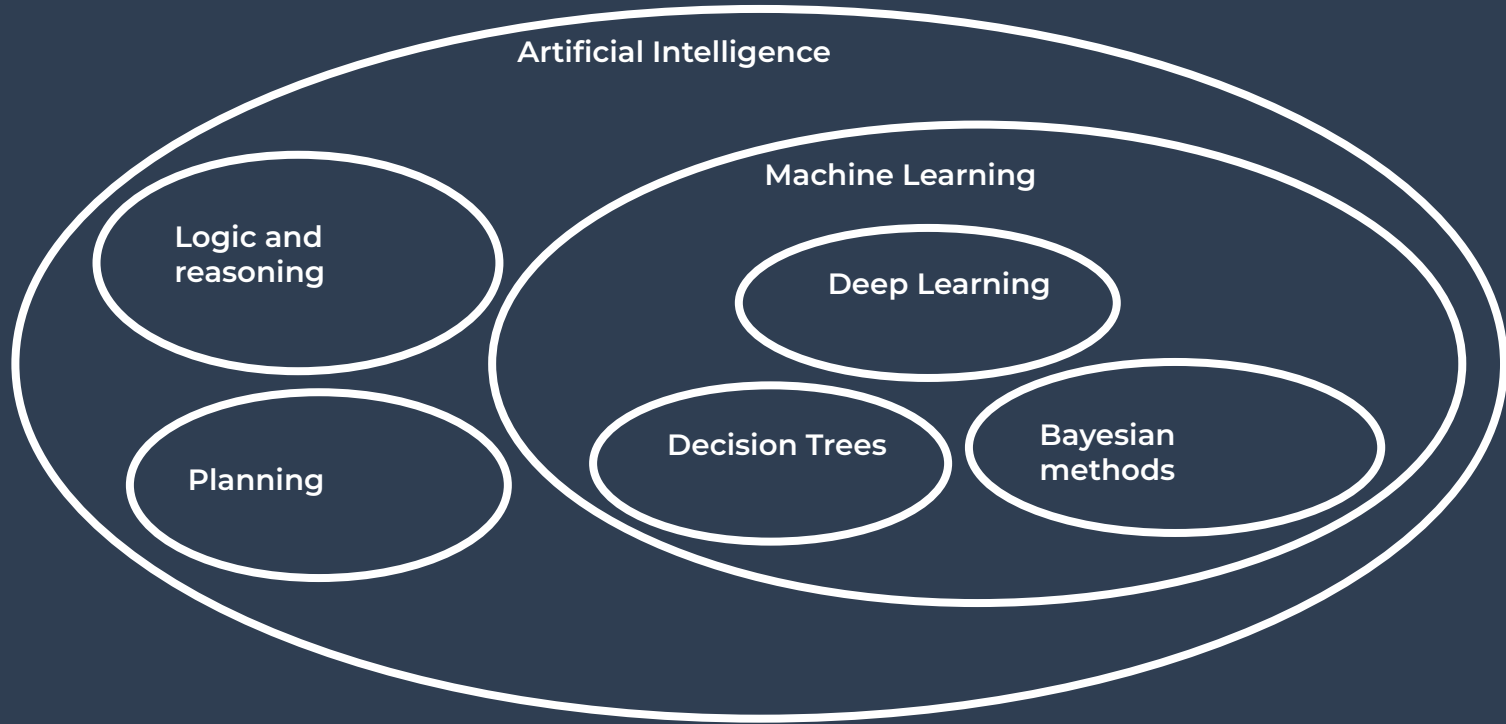
- **Learning from data !**
- Design extracted from data.
- Need time and data to exploit.
- Can be based on decentralized processes (neural network)
- Examples : machine learning, genetic algorithm, optimisation.
- Very efficient if enough data available.

Machine Learning and AI

- **Machine Learning** is a subfield of Artificial Intelligence.
- ML is based on mathematics and statistics.
- **Goal:** machines that can learn from examples.

“ Improving the performance of a program in solving tasks without being explicitly programmed to solve these tasks “

Machine Learning and AI



“ Change behavior in order to become more efficient in the future at a task “

What is learning ?

- Change behavior in order to become more efficient in the future at a task.
- A group of methods for finding and describing structural patterns in data.
- Take observations and model them, describe them, and make predictions.
- Opposed to the algorithmic approach (logic, expert systems) which does not start from data.

What is learning ?


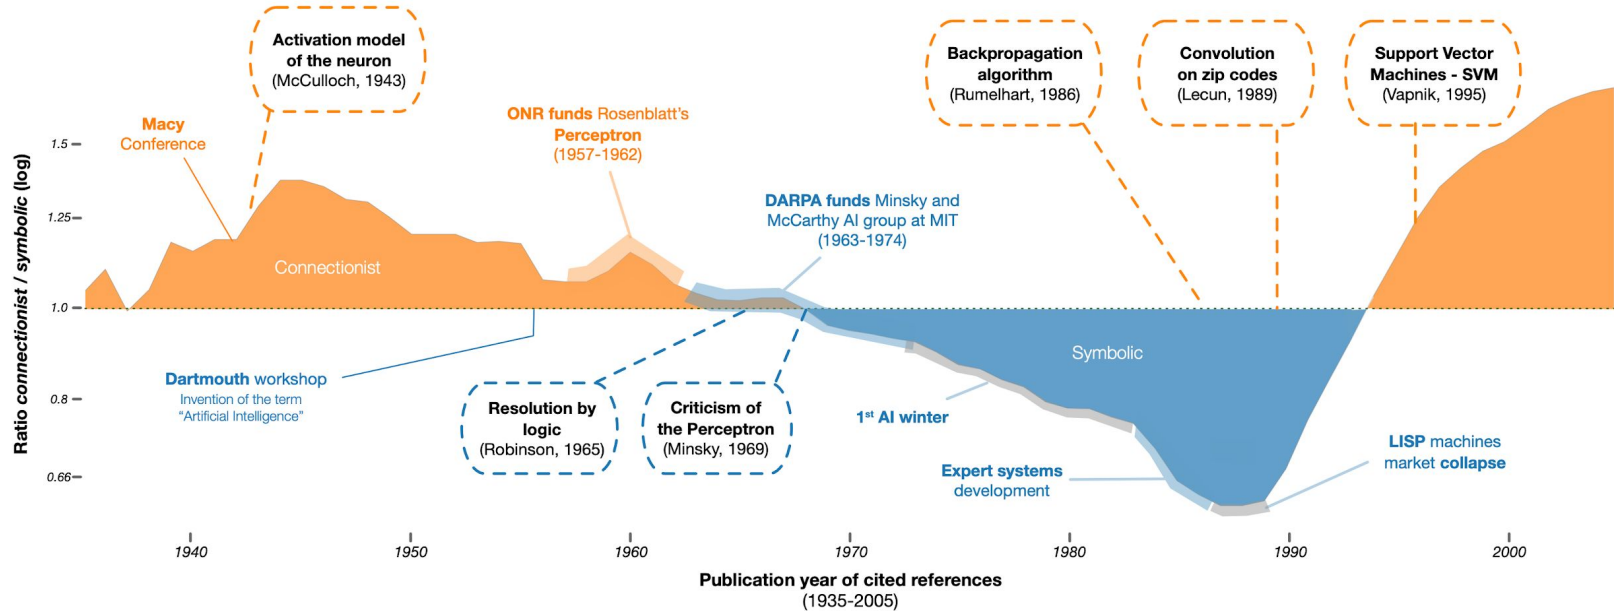
- Find a function $X \rightarrow Y$
- We only have some observations:
 - $(x_1, y_1), (x_2, y_2), (x_3, y_3)$
 - Or learning examples
- We want to predict y for a given x ($y = f(x)$)

What is learning ?

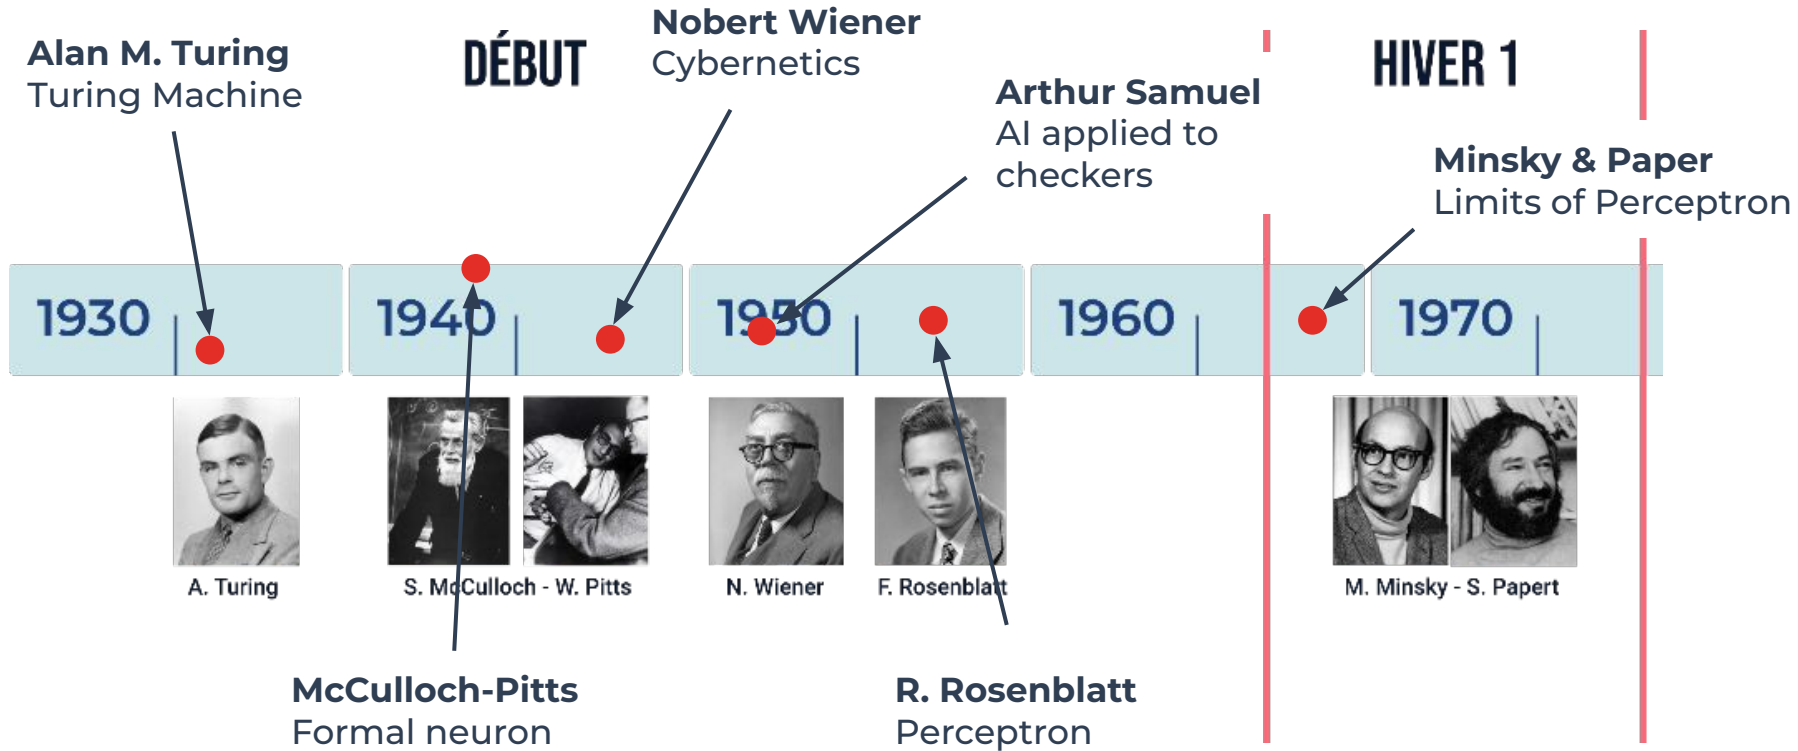
Recall : garbage in, garbage out

Bad, noisy data: nothing to learn.

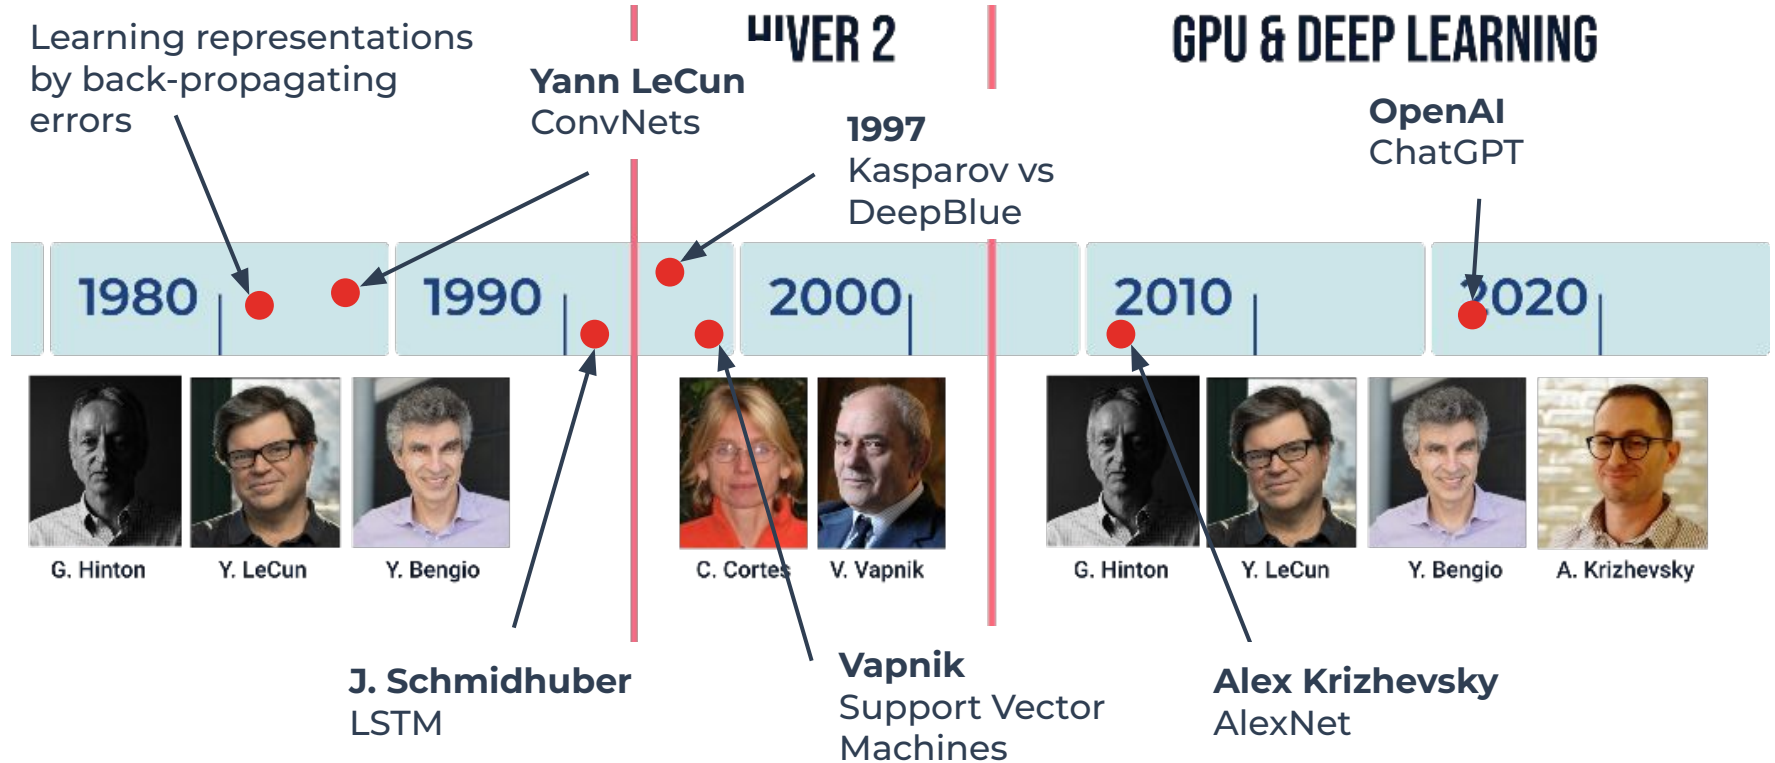
History of Machine Learning



History of Machine Learning



History of Machine Learning




History of Artificial Intelligence

- **On Computable Numbers, with an Application to the Entscheidungsproblem** (Alan M. Turing, 1936)
- **A Logical Calculus of the Ideas Immanent in Nervous Activity** (McCulloch, W. S., Pitts, W., 1943)
- **Cybernetics or Control and Communication in the Animal and the Machine** (Norbert Wiener, 1948)
- **Computing Machinery and Intelligence** (Alan M. Turing, 1950)
- **The Perceptron, A perceiving and recognizing automaton** (R. Rosenblatt, 1957)
- **Perceptrons** (Minsky & Paper, 1969)


History of Artificial Intelligence

- **Learning representations by back-propagating errors** (David E. Rumelhart, Geoffrey E. Hinton & Ronald J. Williams, 1986)
- **Backpropagation Applied to Handwritten Zip Code Recognition** (LeCun, Y. Boser, ..., 1989)

History of Machine Learning



Performance in image recognition




Performance in image recognition

ImageNet classification with Deep Convolutional neural networks

(Alex Krizhevsky, Ilya Sutskever, Geoffrey E. Hinton)

A data-driven world



A data-driven world


The Largest database (2013!)

1. World Data Centre for Climate (6 PB)
2. National Energy Research Computing Center (3 PB)
3. AT&T (312 TB)
4. Google
5. Sprint
6. Choice Point.
7. YouTube (45 TB)
8. Amazon (42 TB)
9. CIA
10. Library of Congress (20 TB)

A data-driven world

- We are overwhelmed by data (amount of data doubled every 20 months)
- We need methods to automatically extract information.
- Data is not information
 - Data record facts in a formal way (0 and 1)
 - Information is the patterns underlying the data.

A GPU-driven world



A GPU-driven world

Trends in GPU price-performance

(Marius Hobbhabn and Tamay Besiroglu)